

Classification

PHYS591000 Spring 2021


Reference: <https://indico.cern.ch/event/619370/> Lecture by Michael Kagan

Outline

- Review: supervised and unsupervised learning
- Binary classification with linear discriminant
- Performance of the classification algorithm: ROC and AUC


Supervised Learning

- Trained with the right answers given
- Typical task: **Classification**




Supervised Learning

- Trained with the right answers given
- Typical task: **Regression**




Unsupervised Learning

- Trained without knowing the right answers
- Typical task: **Clustering**



Unsupervised Learning

- Trained without knowing the right answers
- Typical task: **Dimensionality reduction**




Classification


- Topic for the week
- Focus on **binary classification**: Distinguish signal from background
 - Example: Distinguish particles from hard (head-on) collisions (signal) and particles from soft collisions (background) at the LHC

Binary Classification

- At the LHC we collide bunches of protons



Background



Binary Classification

- Use physics to 'guess' a feature that can separate signals and backgrounds, e.g. energy of a particle
- Expect higher energy (E) for signals from hard collisions


Background

Signal



Binary Classification

- Signals from hard collisions have higher *mean* (average) energy
- There is overlap of signal and background due to *spreads* of their energy distributions



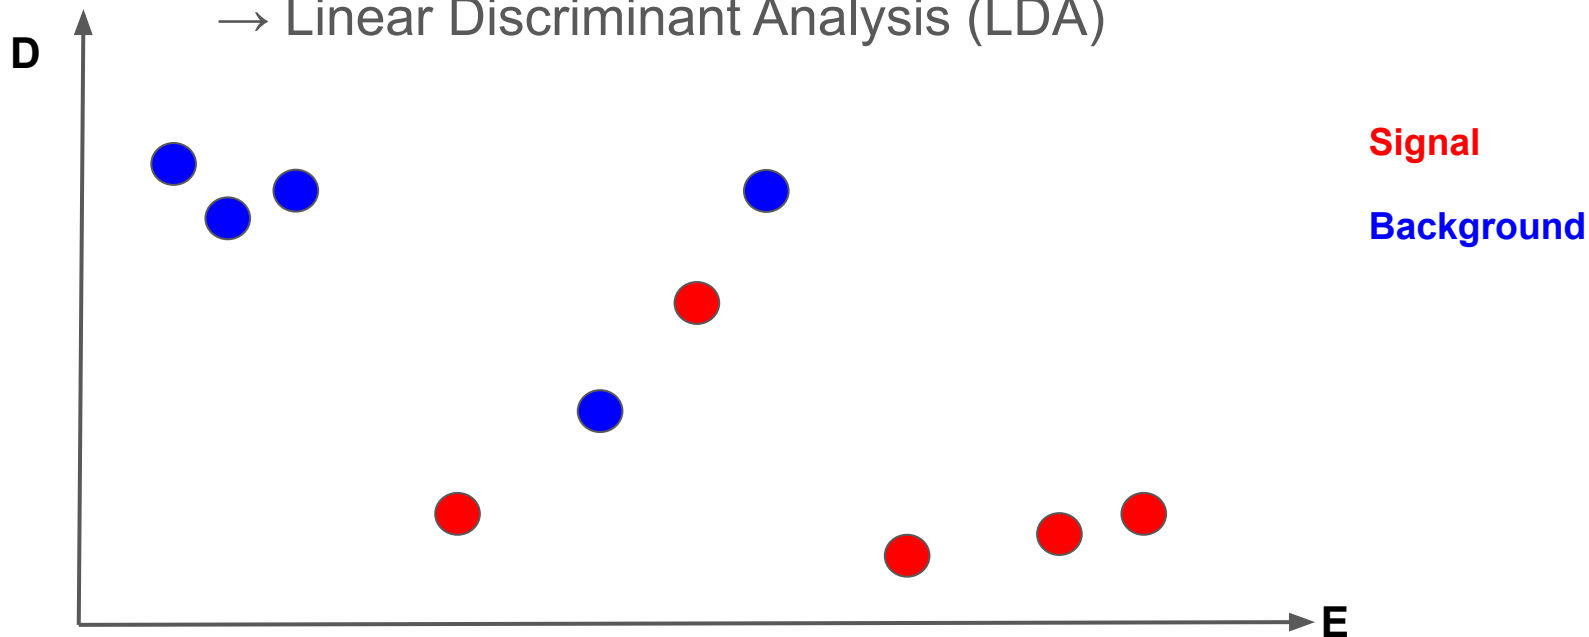
Binary Classification

- Using the energy to separate signals from backgrounds is OK.
- We may be able to do better by using more information, e.g.
D = distance from the primary vertex (main collision point)
-> Expect signals originate from tracks nearer the primary vertex.

Binary Classification


- Want to transform the 2D information into a number (1D)

→ Linear Discriminant Analysis (LDA)



Linear Discriminant Analysis (LDA)

- LDA finds a new axis and projects all data onto the new axis



Linear Discriminant Analysis (LDA)

- The new axis can maximize the separation between the two categories when data are projected on it.


Background

Signal



Linear Discriminant Analysis (LDA)


- The new axis is created by using optimized weights to combine all the information in order to
 - Maximize the distance between the means (μ)
 - Minimize the spreads (“covariance” Σ) within each category



Linear Discriminant Analysis (LDA)

Courtesy of Prof. Kai-Feng Chen (NTU)

- Now let's practice the easiest/simplest algorithm: **Linear discriminant analysis (LDA)**, or even simpler, the **Fisher's discriminant**, by combining the multiple features into one variable:



Calculate the weights (\mathbf{w}_i) to maximize the separation \mathbf{S} .

Linear Discriminant Analysis (LDA)

Courtesy of Prof. Kai-Feng Chen (NTU)

- Consider a set of observables: $\vec{x} = (x_1, x_2, x_3, \dots)$
- For 2 different event classes, the **mean** and **covariance** of the observables are: $\vec{\mu}_0, \vec{\mu}_1, \Sigma_0, \Sigma_1$

$$\vec{\mu} = \langle \vec{x} \rangle \quad \Sigma = \langle (\vec{x} - \vec{\mu}) \cdot (\vec{x} - \vec{\mu})^T \rangle$$

- The separation S is given by
$$S = \frac{(\vec{w} \cdot \vec{\mu}_1 - \vec{w} \cdot \vec{\mu}_0)^2}{\vec{w}^T \Sigma_1 \vec{w} + \vec{w}^T \Sigma_0 \vec{w}}$$

distance of μ
covariance Σ
- The optimal weights can be determined by maximizing the S :

$$\vec{w} \propto (\Sigma_0 + \Sigma_1)^{-1} (\vec{\mu}_1 - \vec{\mu}_0)$$

Linear Discriminant Analysis (LDA)

- It's straightforward to implement the calculation in numpy:

```
mu0 = var0.mean(axis=1)  ← mean values, in shape of (2, )
mu1 = var1.mean(axis=1)
cov0 = np.cov(var0)      ← covariance matrix, in shape of (2, 2)
cov1 = np.cov(var1)

weight = np.dot(linalg.inv(cov1+cov0), mu1-mu0)  ← weight calculation
norm = np.sqrt((weight**2).sum())
weight /= norm
```

Courtesy of Prof. Kai-Feng Chen (NTU)

Linear Discriminant Analysis (LDA)

- Or with Scikit-learn:

```
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
```

```
clf = LinearDiscriminantAnalysis()  
f_train = clf.fit_transform(x_train, y_train) ← “training”
```

Courtesy of Prof. Kai-Feng Chen (NTU)


You'll learn more in the in-class sessions.

Performance of the Classifier

- The LDA maps all the information into one 'score'.
-> Where should we cut on to separate signals from backgrounds?
- The performance of this LDA depends on the cut value we choose.


$X < X_1$: Predicted as Background

$X > X_1$: Predicted as Signal



Performance of the Classifier

- The LDA maps all the information into one 'score'.
-> Where should we cut on to separate signals from background?
- The performance of this LDA depends on the cut value we choose.



Performance of the Classifier

- One way to quantify the performance of the chosen cut is to construct the corresponding **confusion matrix**

	Actual Signal	Actual Background
Predicted as Signal	True Positives	False Positives
Predicted as Background	False Negatives	True Negatives

Performance of the Classifier

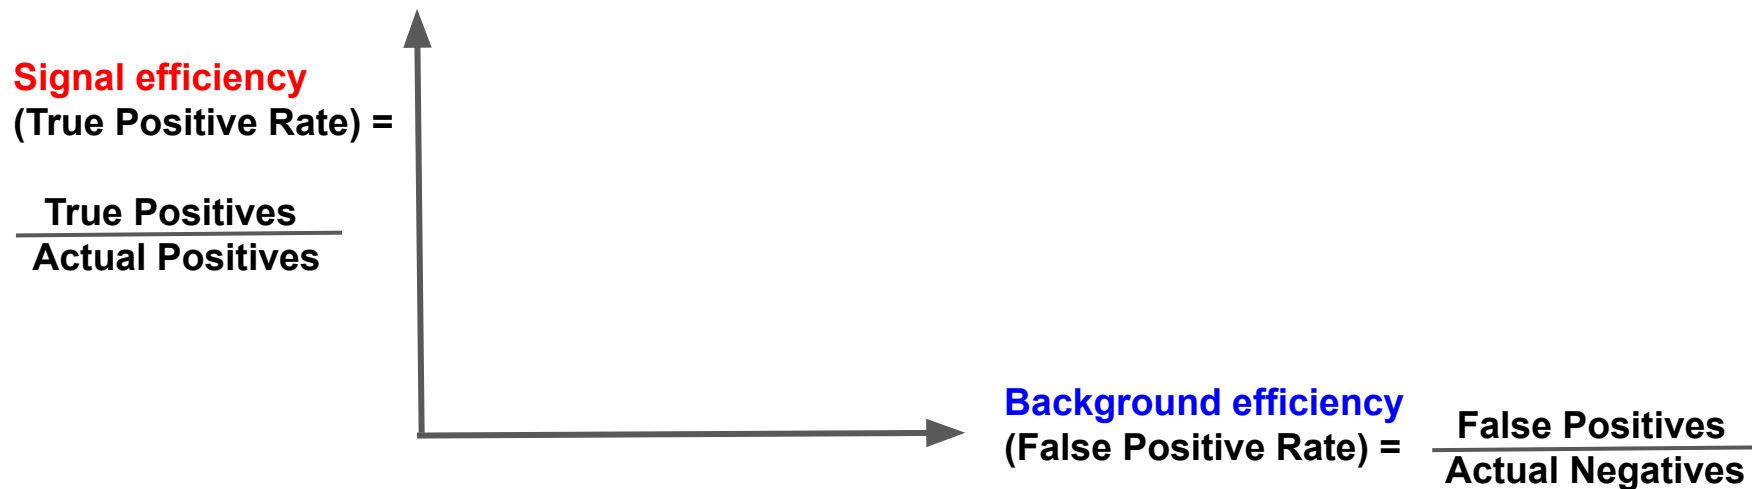
- One way to quantify the performance of the chosen cut is to construct the corresponding **confusion matrix**

X_1	Actual Signal	Actual Background
Predicted as Signal	5	1
Predicted as Background	0	4

X_2	Actual Signal	Actual Background
Predicted as Signal	4	0
Predicted as Background	1	5


Performance of the Classifier

- A good way to summarize all confusion matrices is the **ROC** curve (receiver operating characteristic curve)



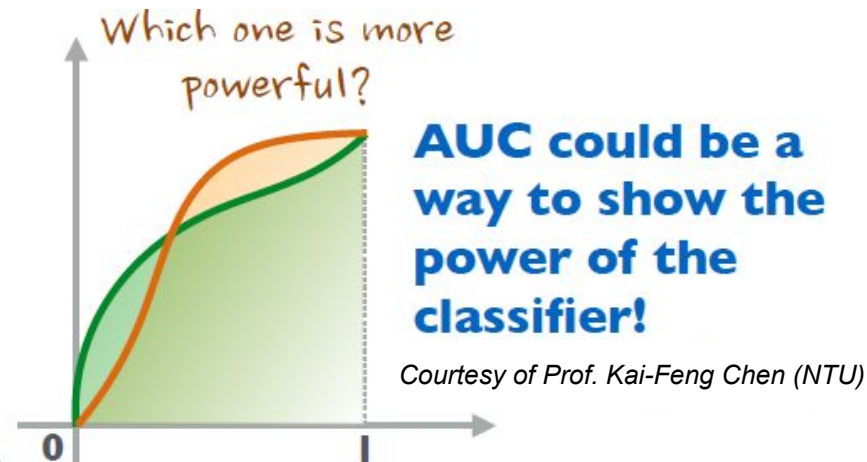
Performance of the Classifier

- The **ROC** curve illustrates the ability of the binary classifier when the discrimination threshold is varied.




Performance of the Classifier

- ROC curves can be used to compare two or more classifiers:
The more it bends away from the diagonal line, the better its performance is.
- Another way to estimate the performance is the **AUC** (area under the curve), which varies from 0.5 (diagonal line) to 1.0 (ideal case)




Summary-I

- The linear discriminant analysis (LDA) algorithm projects all information to a new axis which maximizes the separation of the two categories





Rectangular cuts



Linear discriminant


Summary-II

- The performance of a binary classifier can be estimated by the ROC curve or the AUC




Outlook

- What if we want to do multiclass classification?
- Or if we need a Non-linear Discriminant?



Nonlinear discriminant



Outlook

- What if we want to do multiclass classification?
- Or if we need a Non-linear Discriminant?
- We will make use of more sophisticated algorithms such as support vector machines (SVM), decision trees, neural networks... next time!


LDA in Scikit-learn

- In Scikit-learn LDA makes predictions by estimating the *probability* of an event belongs to each class, assuming the probability of each class is Gaussian and shares the same covariance.
- The predicted class is the one with the highest (posterior) probability.

LDA in Scikit-learn

- The probability is calculated using Bayes' Theorem
 - Then the conditional probability, $P(A|B)$, the probability that an elementary event, known to belong to the set B , and is also a member of set A :

$$P(A \text{ and } B) = P(A|B)P(B) = P(B|A)P(A)$$



Bayes theorem $P(A|B) = P(B|A) \cdot P(A)/P(B)$

Courtesy of Prof. Kai-Feng Chen (NTU)